

What is the Backend Story While Executing Ansible Ad-hoc commands ?

- **Step1:** Ansible reads the list of managed servers from inventory file
- **Step2:** Ansible connects to the managed nodes using SSH with current logged-in user
- **Step3:** Ansible pushes our module to remote nodes (`~/.ansible`) and executes the task on managed node itself using python and it deletes the pushed module once task is completed.
- Ansible uses batches for task execution, which are controlled by a parameter called forks
- The default value for forks is 5, that means ansible executes a task on the first five managed nodes parallelly and waits for the task to complete, and then next five nodes and so on...
- This default forks can be changed using
 - Environment Variable : `export ANSIBLE_FORKS=requiredBatchSize`
 - By Passing `-f requiredBatchSize` in ansible Ad-hoc commands
 - By Keeping default value in `ansible.cfg` with `forks=requiredBatchSize`
- **NOTE:**
 - We must need 1GB per 10 forks + 2GB Reservation for Ansible Controller
 - So, if `forks=20` then we need 2GB and + 2GB Reservation for Ansible Controller , so total 4GB